


CS541 Artificial Intelligence

Guest Lecture on Mean Estimation


Lecturer: Shiwei Zeng

Estimating Average Height

- Assume $D = N(60, 1)$



- Assume $E[D] = 60, \text{Var}[D] = 1$




- Estimator $\hat{\mu} = \frac{1}{n} \sum_{i=1}^n x_i$ $x_i \sim D$ i.i.d.

$$E[\hat{\mu}] = E\left[\frac{1}{n} \sum_{i=1}^n x_i\right] = \frac{1}{n} \sum_{i=1}^n E[x_i] = \mu$$

$$\text{Var}[\hat{\mu}] = \text{Var}\left[\frac{1}{n} \sum_{i=1}^n x_i\right] = \frac{1}{n^2} \sum_{i=1}^n \underbrace{\text{Var}[x_i]} = \frac{1}{n}$$



ME in Higher Dimension



D

$$E[D] = ?$$

When Data is Noisy

$O(\cdot)$

$\xi < \frac{1}{2}$ $\xi = 3\%$
 $\xi \cdot C \quad C > 0$

- 1-dimensional: (a lower bound)

$\|\hat{M} - M\|_2^2 \geq \Omega(\xi)$ if corrupt ξ -fraction.

$D_1 = N(\mu_1, 1) \quad D_2 = N(\mu_2, 1) \quad \text{construct } Q_1 \text{ and } Q_2 \text{ s.t.}$

$\|\mu_1 - \mu_2\| = \Omega(\xi) \text{ and}$

$\boxed{D_\xi} = (1 - \xi) D_1 + \xi Q_1 = (1 - \xi) D_2 + \xi Q_2 \quad (1)$

Let ϕ_1 be pdf of D_1 , ϕ_2 pdf of D_2 . Let μ_1, μ_2 be s.t. the total variance distance between D_1, D_2 is

$\frac{1}{2} \int |\phi_1 - \phi_2| dx = \frac{\xi}{1 - \xi} \implies \|\mu_1 - \mu_2\| \geq \frac{2\xi}{1 - \xi}.$

$Q_1: \frac{1 - \xi}{\xi} (\phi_2 - \phi_1) \cdot \mathbb{1}_{\phi_2 > \phi_1} \quad \text{and} \quad Q_2: \frac{1 - \xi}{\xi} (\phi_1 - \phi_2) \cdot \mathbb{1}_{\phi_1 > \phi_2}.$

$$D_1 = N(\mu_1, 1) \quad D_2 = N(\mu_2, 1)$$

$$\|\mu_1 - \mu_2\| = \Omega(\varepsilon) \text{ and}$$

$$\boxed{D_\varepsilon} = (1-\varepsilon) D_1 + \varepsilon Q_1 = (1-\varepsilon) D_2 + \varepsilon Q_2 \quad (1)$$

Let ϕ_1 be pdf of D_1 , ϕ_2 pdf of D_2 . Let μ_1, μ_2 be s.t. the total variate distance between D_1, D_2 is

$$\frac{1}{2} \int |\phi_1 - \phi_2| dx = \frac{\varepsilon}{1-\varepsilon} \implies \|\mu_1 - \mu_2\| \geq \frac{2\varepsilon}{1-\varepsilon} \geq \Omega(\varepsilon)$$

$\frac{2\varepsilon}{1-\varepsilon} \geq 2\varepsilon$
 $\in (\frac{1}{2}, 1)$

$$Q_1: \frac{1-\varepsilon}{\varepsilon} (\phi_2 - \phi_1) \cdot \mathbb{1}_{\phi_2 > \phi_1} \text{ and } Q_2: \frac{1-\varepsilon}{\varepsilon} (\phi_1 - \phi_2) \mathbb{1}_{\phi_1 > \phi_2}.$$

$$\begin{aligned} & (1-\varepsilon) \phi_1 + \varepsilon \cdot \frac{1-\varepsilon}{\varepsilon} (\phi_2 - \phi_1) \cdot \mathbb{1}_{\phi_2 > \phi_1} \\ &= (1-\varepsilon) \phi_1 + (1-\varepsilon) (\phi_2 - \phi_1) \cdot \mathbb{1}_{\phi_2 > \phi_1} \end{aligned}$$

$$\mathbb{1}_{\phi_2 > \phi_1} = \begin{cases} 1 & \phi_2 > \phi_1 \\ 0 & \phi_2 < \phi_1 \end{cases}$$


$$= \begin{cases} (1-\varepsilon) \cdot \phi_2 & \phi_1 \leq \phi_2 \\ (1-\varepsilon) \cdot \phi_1 & \phi_1 > \phi_2 \end{cases}$$

$$(1-\varepsilon) \phi_2 + \varepsilon \cdot \frac{1-\varepsilon}{\varepsilon} (\phi_1 - \phi_2) \mathbb{1}_{\phi_1 > \phi_2}$$

$$= \begin{cases} (1-\varepsilon) \cdot \phi_1 & \phi_1 > \phi_2 \\ (1-\varepsilon) \cdot \phi_2 & \phi_1 < \phi_2 \end{cases} \quad \phi_1 = \phi_2$$

Robust Mean Estimation


Mean Estimation



D

$$E[D] = ?$$

ϵ -robust Mean Estimation



$D + D'$

$$E[D] = ?$$

Natural approaches

- Learn each coordinate separately

$$\begin{aligned} \hat{\mu} \text{ s.t. } & \|\hat{\mu}_i - \mu_i\| \geq \Omega(\varepsilon). \\ \text{in } n\text{-dim. } & \|\hat{\mu} - \mu\|_2^2 \geq \sum_{i=1}^n (\hat{\mu}_i - \mu_i)^2 = n \cdot \varepsilon^2 \\ & \|\hat{\mu} - \mu\|_2 \geq \sqrt{n} \cdot \varepsilon \end{aligned}$$
$$\begin{aligned} n &= 10^6 \\ \sqrt{n} \cdot \varepsilon &= 10^3 \cdot \varepsilon \leq (0, \tfrac{1}{2}) \end{aligned}$$

Natural approaches

e^n

- Maximum Likelihood Estimator

Negative Log Likelihood

$$\min NLL(F, x_1, \dots, x_m) = - \sum_{i=1}^m \log F(x_i)$$

$$S = \{x_i\}_{i=1}^m$$

Assume F is Gaussian: $F(x_i) = \frac{1}{\sqrt{2\pi}} \cdot e^{-\frac{\|x_i - \mu\|_2^2}{2}}$

$$\min NLL = \min - \sum_{i=1}^m \log \frac{1}{\sqrt{2\pi}} \cdot e^{-\frac{\|x_i - \mu\|_2^2}{2}}$$

$$= \min_{\mu} \left(-N \cdot \log \frac{1}{\sqrt{2\pi}} + \sum_{i=1}^m \frac{\|x_i - \mu\|_2^2}{2} \right)$$

$$\Rightarrow \min_{\mu} \frac{1}{2} \sum_{i=1}^m \|x_i - \mu\|_2^2$$

$\hat{\mu}$: empirical mean
 $= \frac{1}{m} \sum_{i=1}^m x_i$

Efficient Algorithm – Convex Programming

~ 7:30

Output a $\hat{w} = (w_1, w_2, \dots, w_m)$ such that

$$\hat{\mu} = \frac{1}{m} \sum_{i=1}^m w_i x_i \quad \text{is close to } \mu.$$

by solving a convex program.

$$O(n^b)$$

$$D = N(\underline{\mu}, I) \quad \forall v \in \mathbb{R}^n \quad \Pr_{x \sim D}[|v \cdot (x - \underline{\mu})| \geq t] \leq \exp(-\frac{t^2}{2}) \quad \star \quad \mathcal{E}(0, \frac{1}{2}) \quad \tau = 0.01 \quad 0.99$$

Efficient Robust Mean Estimation - Filter

T : corrupted dataset.

$$\Pr_{x \sim D}[|p(x) - E[p(x)]| \geq t] \leq \boxed{\frac{1}{m} \sum_{i=1}^m x_i \quad \frac{1}{m} \sum_{i=1}^m (x_i - \mu_T)(x_i - \mu_T)^T}$$

1. Compute empirical mean and covariance $\mu_T, \Sigma_T = I$
2. Compute largest eigenvalue λ^* of $\Sigma_T - I$, and eigenvector v^*
 $= 0$
3. If λ^* is small, return μ_T

4. Otherwise, find $t > C_1$ such that

$$C_1, C_2, C_3 > 0$$

$$\Pr_{X \in T}[|v^* \cdot (X - \mu_T)| > t] > C_2 e^{-t^2/2} + \frac{C_3 \varepsilon}{t^2 \log(n \log \frac{n}{\varepsilon \tau})}$$


$= \delta$ dependent on m, τ

5. Remove X such that $|v^* \cdot (X - \mu_T)| > t$, go back to step 1

\sqrt{n} , n -dim

$$\hat{\mu} \rightarrow \mu$$


$$\|\hat{\mu} - \mu\|_2 \leq o(\varepsilon) \leq 1$$



List-decodable Mean Estimation

T : corrupted data set.

Mean Estimation




D

$$E[D] = ?$$

List-Decodable Mean Estimation

α : fraction of clean samples.



k golden queries

\rightarrow a list of e^k estimated means.

$$E[D] \in \{\mu_1, \dots, \mu_m\}$$

$\log(m)$ queries

$$\rightarrow m = O\left(\frac{1}{\alpha}\right)$$


T_i is α -good: α -fraction of T_i are clean.

Algorithm: Multi-filtering

- A tree of subsets T_i 's,
- Iterate through each node
 - (1) **Create** a leaf node, an **estimate** $\hat{\mu}_i \leftarrow$
 - (2) **Create** child nodes, **subsets** T_i 's
 - a. One node, cleaner set
 - b. Two nodes, overlapping subsets
 - (3) **Delete** if it **can't be** α -good.
- No more filtering, then **return** all $\hat{\mu}_i$'s

$$\min_{i \in \{1, \dots, m\}} \|\hat{\mu}_i - \mu\|_2 = O\left(\frac{1}{\sqrt{\alpha}}\right) = 1000$$

$\alpha = 10^{-6}$



α -good T_i = α -fraction of T_i .

$$\hat{\mu} \leftarrow \mu \quad O\left(\frac{1}{\alpha^{\frac{1}{2d}}}\right) \Rightarrow O\left(\frac{1}{\alpha^{\frac{1}{2d}}}\right) \quad d: \text{degree}$$